

The IPython Notebook: open, reproducible scientific computing

Matthias Bussonnier¹, Jonathan Frederic², Bradley M. Froehle³, Brian E. Granger², Paul Ivanov³, Thomas Kluyver³, Fernando Perez³, Benjamin Ragan-Kelley³ and Zachary Sailer²

¹Affiliation of Matthias

²Cal Poly State University

³University of CA, Berkeley

June 30, 2013

Abstract

While computing has become a foundation of all it is challenging for researchers .

1 Introduction

2 The lifecycle of research

3 The IPython Notebook

3.1 Web application

3.2 Notebook document format

3.3 Installation

4 Collaboration

5 Broader ecosystem

6 Future directions

Dexy Snippets

The notebook `notebooks/Chapter1_Introduction.ipynb`— uses format 3.0.

Markdown Content

Here is an example of a markdown cell:

3What is the mean of λ_1 **given** we know τ is less than 45. That is, suppose we have new information as we know for certain that the change in behaviour occurred before day 45. What is the expected value of λ_1 now? (You do not need to redo the PyMC part, just consider all instances where ‘ $\tau_{samples} < 45$ ’.)

Here is the same cell in a Verbatim block:

```
3\ . What is the mean of  $\lambda_1$  given we know  $\tau$  is less than 45. That is,
```

Here is the same cell presented as highlighted markup:

```
3\ . What is the mean of  $\lambda_1$  given we know  $\tau$  is less than 45. That is,
```

Python Content

```
"""
```

The book uses a custom matplotlibrc file, which provides the unique styles for matplotlib. If executing this book, and you wish to use the book's styling, provided are two options:

1. Overwrite your own matplotlibrc file with the rc-file provided in the book's style. See <http://matplotlib.org/users/customizing.html>
2. Also in the styles is `bmh_matplotlibrc.json` file. This can be used to update the style in only this notebook. Try running the following code:

```
import json
s = json.load( open("../styles/bmh_matplotlibrc.json") )
matplotlib.rcParams.update(s)
```

```
"""
```

```
#the code below can be passed over, as it is currently not important.
```

```
%pylab inline
```

```
figsize( 11, 9)
```

```
import scipy.stats as stats
```

```
dist = stats.beta
```

```
n_trials = [0,1,2,3,4,5,8,15, 50, 500]
```

```
data = stats.bernoulli.rvs(0.5, size = n_trials[-1] )
```

```
x = np.linspace(0,1,100)
```

```
for k, N in enumerate(n_trials):
```

```
    sx = subplot( len(n_trials)/2, 2, k+1)
```

```
    plt.xlabel("$p$, probability of heads") if k in [0,len(n_trials)-1] else None
```

```
    plt.setp(sx.get_yticklabels(), visible=False)
```

```

heads = data[:N].sum()
y = dist.pdf(x, 1 + heads, 1 + N - heads )
plt.plot( x, y, label= "observe %d tosses,\n %d heads"%(N,heads) )
plt.fill_between( x, 0, y, color="#348ABD", alpha = 0.4 )
plt.vlines( 0.5, 0, 4, color = "k", linestyle = "--", lw=1 )

leg = plt.legend()
leg.get_frame().set_alpha(0.4)
plt.autoscale(tight = True)

plt.suptitle( "Bayesian updating of posterior probabilities",
             y = 1.02,
             fontsize = 14);

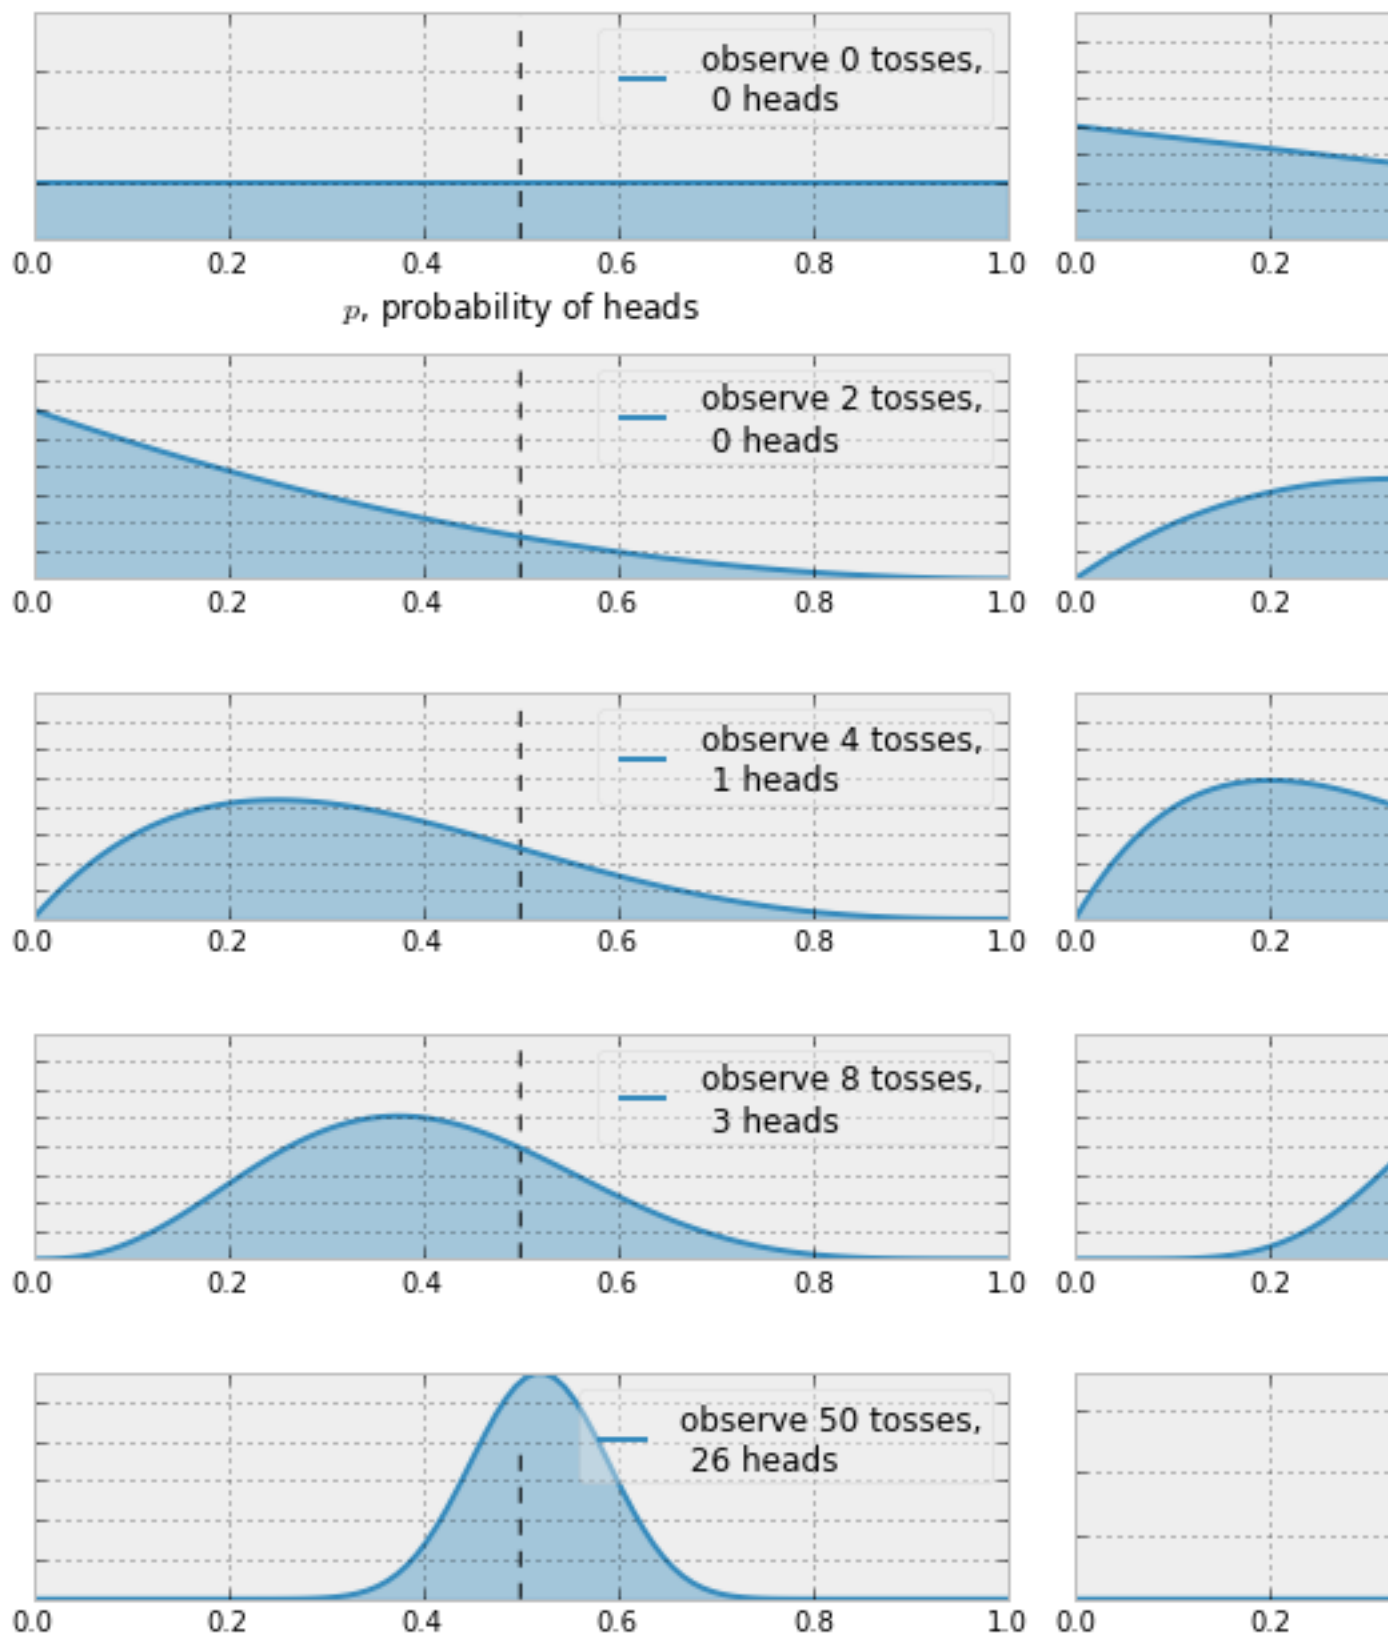
plt.tight_layout()

```

Welcome to pylab, a matplotlib-based Python environment [backend: module://IPython.zmq.p

For more information, type 'help(pylab)'.

Bayesian updating of posterior probability



Iterating over Cells and Documents

This notebook has 47 cells and 59 documents based on these cells.

Cells

Here is a list of cells:

- notebooks/Chapter1_Introduction--0
- notebooks/Chapter1_Introduction--1
- notebooks/Chapter1_Introduction--2
- notebooks/Chapter1_Introduction--3
- notebooks/Chapter1_Introduction--4
- notebooks/Chapter1_Introduction--5
- notebooks/Chapter1_Introduction--6
- notebooks/Chapter1_Introduction--7
- notebooks/Chapter1_Introduction--8
- notebooks/Chapter1_Introduction--9
- notebooks/Chapter1_Introduction--10
- notebooks/Chapter1_Introduction--11
- notebooks/Chapter1_Introduction--12
- notebooks/Chapter1_Introduction--13
- notebooks/Chapter1_Introduction--14
- notebooks/Chapter1_Introduction--15
- notebooks/Chapter1_Introduction--16
- notebooks/Chapter1_Introduction--17
- notebooks/Chapter1_Introduction--18
- notebooks/Chapter1_Introduction--19
- notebooks/Chapter1_Introduction--20
- notebooks/Chapter1_Introduction--21
- notebooks/Chapter1_Introduction--22

- notebooks/Chapter1_Introduction--23
- notebooks/Chapter1_Introduction--24
- notebooks/Chapter1_Introduction--25
- notebooks/Chapter1_Introduction--26
- notebooks/Chapter1_Introduction--27
- notebooks/Chapter1_Introduction--28
- notebooks/Chapter1_Introduction--29
- notebooks/Chapter1_Introduction--30
- notebooks/Chapter1_Introduction--31
- notebooks/Chapter1_Introduction--32
- notebooks/Chapter1_Introduction--33
- notebooks/Chapter1_Introduction--34
- notebooks/Chapter1_Introduction--35
- notebooks/Chapter1_Introduction--36
- notebooks/Chapter1_Introduction--37
- notebooks/Chapter1_Introduction--38
- notebooks/Chapter1_Introduction--39
- notebooks/Chapter1_Introduction--40
- notebooks/Chapter1_Introduction--41
- notebooks/Chapter1_Introduction--42
- notebooks/Chapter1_Introduction--43
- notebooks/Chapter1_Introduction--44
- notebooks/Chapter1_Introduction--45
- notebooks/Chapter1_Introduction--46

Documents

Here is a list of documents:

- notebooks/Chapter1_Introduction--0.md
- notebooks/Chapter1_Introduction--1.md
- notebooks/Chapter1_Introduction--2.md
- notebooks/Chapter1_Introduction--3.md
- notebooks/Chapter1_Introduction--4.md
- notebooks/Chapter1_Introduction--5.md
- notebooks/Chapter1_Introduction--6.md
- notebooks/Chapter1_Introduction--7-input.py
- notebooks/Chapter1_Introduction--7-output-0.txt
- notebooks/Chapter1_Introduction--7-output-1.png
- notebooks/Chapter1_Introduction--8.md
- notebooks/Chapter1_Introduction--9.md
- notebooks/Chapter1_Introduction--10.md
- notebooks/Chapter1_Introduction--11.md
- notebooks/Chapter1_Introduction--12-input.py
- notebooks/Chapter1_Introduction--12-output-0.png
- notebooks/Chapter1_Introduction--13.md
- notebooks/Chapter1_Introduction--14-input.py
- notebooks/Chapter1_Introduction--14-output-0.png
- notebooks/Chapter1_Introduction--15.md
- notebooks/Chapter1_Introduction--16.md
- notebooks/Chapter1_Introduction--17-input.py
- notebooks/Chapter1_Introduction--17-output-0.png
- notebooks/Chapter1_Introduction--18.md
- notebooks/Chapter1_Introduction--19-input.py

- notebooks/Chapter1_Introduction--19-output-0.png
- notebooks/Chapter1_Introduction--20.md
- notebooks/Chapter1_Introduction--21.md
- notebooks/Chapter1_Introduction--22-input.py
- notebooks/Chapter1_Introduction--22-output-0.png
- notebooks/Chapter1_Introduction--23.md
- notebooks/Chapter1_Introduction--24-input.py
- notebooks/Chapter1_Introduction--25.md
- notebooks/Chapter1_Introduction--26-input.py
- notebooks/Chapter1_Introduction--26-output-0.txt
- notebooks/Chapter1_Introduction--27-input.py
- notebooks/Chapter1_Introduction--28.md
- notebooks/Chapter1_Introduction--29-input.py
- notebooks/Chapter1_Introduction--30.md
- notebooks/Chapter1_Introduction--31-input.py
- notebooks/Chapter1_Introduction--31-output-0.txt
- notebooks/Chapter1_Introduction--31-output-1.txt
- notebooks/Chapter1_Introduction--32-input.py
- notebooks/Chapter1_Introduction--33-input.py
- notebooks/Chapter1_Introduction--33-output-0.png
- notebooks/Chapter1_Introduction--34.md
- notebooks/Chapter1_Introduction--35.md
- notebooks/Chapter1_Introduction--36-input.py
- notebooks/Chapter1_Introduction--36-output-0.png
- notebooks/Chapter1_Introduction--37.md
- notebooks/Chapter1_Introduction--38.md
- notebooks/Chapter1_Introduction--39-input.py

- notebooks/Chapter1_Introduction--40.md
- notebooks/Chapter1_Introduction--41-input.py
- notebooks/Chapter1_Introduction--42.md
- notebooks/Chapter1_Introduction--43-input.py
- notebooks/Chapter1_Introduction--44.md
- notebooks/Chapter1_Introduction--45-input.py
- notebooks/Chapter1_Introduction--46-input.py

Here are the contents of documents:

Here are the contents of notebooks/Chapter1_Introduction--0.md:

Probabilistic Programming

=====

and Bayesian Methods **for** Hackers

=====

#####Version 0.1

Welcome to *Bayesian Methods **for** Hackers*. The full Github repository, and additional ch

Here are the contents of notebooks/Chapter1_Introduction--1.md:

Chapter 1

=====

Here are the contents of notebooks/Chapter1_Introduction--2.md:

The Philosophy of Bayesian Inference

> You are a skilled programmer, but bugs still slip into your code. After a particularly

If you think **this** way, then congratulations, you already are a Bayesian practitioner! Ba

Here are the contents of notebooks/Chapter1_Introduction--3.md:

###The Bayesian state of mind

Bayesian inference differs from more traditional statistical inference by preserving *un

The Bayesian world-view interprets probability **as** measure of *believability **in** an event*

For **this** to be clearer, we consider an alternative interpretation of probability: *Frequentists, on the other hand, have a more intuitive approach. Bayesians interpret a probability as a belief measure. Notice **in** the paragraph above, I assigned the belief (probability) measure to an individual. Here are some examples:

- I flip a coin, and we both guess the result. We would both agree, assuming the coin **is** fair.
- Your code either has a bug **in** it or not, but we **do** not know **for** certain which **is** true.
- A medical patient **is** exhibiting symptoms \$x\$, \$y\$ and \$z\$. There are a number of diseases that could cause these symptoms.

This philosophy of treating beliefs **as** probability **is** natural to humans. We employ it consistently. To align ourselves **with** traditional probability notation, we denote our belief about event A given evidence X as $P(A | X)$. John Maynard Keynes, a great economist and thinker, said **"When the facts change, I change my mind."**

1\ . $P(A)$: \;\;\\$ the coin has a **50** percent chance of being heads. $P(A | X)$: \;\;\\$ You look at the coin and see heads.

2\ . $P(A)$: \;\;\\$ This big, complex code likely has a bug **in** it. $P(A | X)$: \;\;\\$ The code runs without error.

3\ . $P(A)$: \;\;\\$ The patient could have any number of diseases. $P(A | X)$: \;\;\\$ Perform a medical test.

It's clear that **in** each example we did not completely discard the prior belief after seeing the evidence. By introducing prior uncertainty about events, we are already admitting that any guess will be wrong.

Here are the contents of notebooks/Chapter1_Introduction--4.md:

```
###Bayesian Inference in Practice
```

If frequentist and Bayesian inference were programming functions, **with** inputs being statements and **returning** probabilities, we could write them as follows. For example, **in** our debugging problem above, calling the frequentist **function with** the answer to the question "Is there a bug in the code?" would return the probability of the code being correct.

```
> *YES*, with probability 0.8; *NO*, with probability 0.2
```

This **is** very different from the answer the frequentist **function** returned. Notice that the Bayesian **function** returns the probability of the code being correct, while the frequentist **function** returns the probability of the code being incorrect.

####Incorporating evidence

As we acquire more and more instances of evidence, our prior belief *is* *washed out* by t

Denote N as the number of instances of evidence we possess. As we gather an *infinite*

One may think that *for* large N , one can be indifferent between the two techniques sinc

> Sample sizes are never large. If N *is* too small to get a sufficiently-precise estima

Are frequentist methods incorrect then?

****No.****

Frequentist methods are still useful or state-of-the-art *in* many areas. Tools like Least

A note on *Big Data*

Paradoxically, big data's predictive analytic problems are actually solved by relatively

The much more difficult analytic problems involve *medium data* and, especially troubles

Here are the contents of notebooks/Chapter1_Introduction--5.md:

Our Bayesian framework

We are interested *in* beliefs, which can be interpreted *as* probabilities by thinking Baye

Secondly, we observe our evidence. To *continue* our buggy-code example: *if* our code passe

```
\begin{align}
P( A | X ) &= \frac{ P(X | A) P(A) }{ P(X) } \\\[5pt]
&\propto P(X | A) P(A)\;;\; (\propto \text{is proportional to } )
\end{align}
```

The above formula *is* not unique to Bayesian inference: it *is* a mathematical fact *with* us

Here are the contents of notebooks/Chapter1_Introduction--6.md:

Example: Mandatory coin-flip example

Every statistics text must contain a coin-flipping example, I'll use it here to get it o

We begin to flip a coin, and record the observations: either $\$H\$$ or $\$T\$$. This `is` our obs

Below we plot a sequence of updating posterior probabilities `as` we observe increasing am

Here are the contents of `notebooks/Chapter1_Introduction--7-input.py`:

```
"""
```

The book uses a custom matplotlibrc file, which provides the unique styles for matplotlib. If executing this book, and you wish to use the book's styling, provided are two options:

1. Overwrite your own matplotlibrc file with the rc-file provided in the book's style. See <http://matplotlib.org/users/customizing.html>
2. Also in the styles is `bmh_matplotlibrc.json` file. This can be used to update the style in only this notebook. Try running the following code:

```
import json
s = json.load( open("../styles/bmh_matplotlibrc.json") )
matplotlib.rcParams.update(s)
```

```
"""
```

```
#the code below can be passed over, as it is currently not important.
```

```
%pylab inline
```

```
figsize( 11, 9)
```

```
import scipy.stats as stats
```

```
dist = stats.beta
```

```
n_trials = [0,1,2,3,4,5,8,15, 50, 500]
```

```
data = stats.bernoulli.rvs(0.5, size = n_trials[-1] )
```

```
x = np.linspace(0,1,100)
```

```
for k, N in enumerate(n_trials):
```

```
    sx = subplot( len(n_trials)/2, 2, k+1)
```

```
    plt.xlabel(" $p$ , probability of heads") if k in [0,len(n_trials)-1] else None
```

```
    plt.setp(sx.get_yticklabels(), visible=False)
```

```
    heads = data[:N].sum()
```

```
    y = dist.pdf(x, 1 + heads, 1 + N - heads )
```

```
    plt.plot( x, y, label= "observe %d tosses,\n %d heads"%(N,heads) )
```

```
    plt.fill_between( x, 0, y, color="#348ABD", alpha = 0.4 )
```

```
    plt.vlines( 0.5, 0, 4, color = "k", linestyle = "--", lw=1 )
```

```
    leg = plt.legend()
```

```
    leg.get_frame().set_alpha(0.4)
```

```
    plt.autoscale(tight = True)
```

```
plt.suptitle( "Bayesian updating of posterior probabilities",
             y = 1.02,
             fontsize = 14);
```

```
plt.tight_layout()
```

Here are the contents of notebooks/Chapter1_Introduction--7-output-0.txt:
 Here are the contents of notebooks/Chapter1_Introduction--7-output-1.png:
 Here are the contents of notebooks/Chapter1_Introduction--8.md:

The posterior probabilities are represented by the curves, and our confidence *is* proportionate to the area under the curves.

Notice that the plots are not always *peaked* at 0.5. There *is* no reason it should be: remember that the prior is uniform.

The next example *is* a simple demonstration of the mathematics of Bayesian inference.

Here are the contents of notebooks/Chapter1_Introduction--9.md:

#####Example: Bug, or just sweet, unintended feature?

Let A denote the event that our code has *no bugs* in it. Let X denote the event that the code passes the tests.

We are interested in $P(A|X)$, i.e. the probability of no bugs, given our debugging test results.

What *is* $P(X | A)$, i.e., the probability that the code passes X tests *given* there are no bugs?

$P(X)$ *is* a little bit trickier: The event X can be divided into two possibilities, even if there are no bugs.

Here are the contents of notebooks/Chapter1_Introduction--10.md:

```
\begin{align}
P(X) &= P(X \text{ and } A) + P(X \text{ and } \sim A) \\\[5pt]
&= P(X|A)P(A) + P(X | \sim A)P(\sim A)\\[5pt]
&= P(X|A)p + P(X | \sim A)(1-p)
\end{align}
```

Here are the contents of notebooks/Chapter1_Introduction--11.md:

We have already computed $P(X|A)$ above. On the other hand, $P(X | \sim A)$ *is* subjective.

```
\begin{align}
P(A | X) &= \frac{1 \cdot p}{1 \cdot p + 0.5 (1-p)} \\\[5pt]
&= \frac{2p}{1+p}
\end{align}
```

This *is* the posterior probability. What does it look like *as* a *function* of our prior, p ?

Here are the contents of notebooks/Chapter1_Introduction--12-input.py:

```
figsize(12.5,4)
p = np.linspace( 0,1, 50)
plt.plot( p, 2*p/(1+p), color = "#348ABD", lw = 3 )
#plt.fill_between( p, 2*p/(1+p), alpha = .5, facecolor = ["#A60628"])
plt.scatter( 0.2, 2*(0.2)/1.2, s = 140, c = "#348ABD" )
plt.xlim( 0, 1)
plt.ylim( 0, 1)
plt.xlabel( "Prior,  $P(A) = p$ " )
plt.ylabel( "Posterior,  $P(A|X)$ , with  $P(A) = p$ " )
plt.title( "Are there bugs in my code?" );
```

Here are the contents of notebooks/Chapter1_Introduction--12-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--13.md:

We can see the biggest gains if we observe the X tests passed when the prior probability

Recall that the prior is a probability: p is the prior probability that there are no

Similarly, our posterior is also a probability, with $P(A | X)$ the probability there is

Here are the contents of notebooks/Chapter1_Introduction--14-input.py:

```
figsize( 12.5, 4 )
colours = ["#348ABD", "#A60628"]

prior = [0.20, 0.80]
posterior = [1./3, 2./3]
plt.bar( [0,.7], prior ,alpha = 0.70, width = 0.25, \
        color = colours[0], label = "prior distribution",
        lw = "3", edgecolor = colours[0])

plt.bar( [0+0.25,.7+0.25], posterior ,alpha = 0.7, \
        width = 0.25, color = colours[1],
        label = "posterior distribution",
        lw = "3", edgecolor = colours[1])

plt.xticks( [0.20,.95], ["Bugs Absent", "Bugs Present"] )
plt.title("Prior and Posterior probability of bugs present, prior = 0.2")
plt.ylabel("Probability")
plt.legend(loc="upper left");
```

Here are the contents of notebooks/Chapter1_Introduction--14-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--15.md:

Notice that after we observed X occur, the probability of bugs being absent increased. This was a very simple example of Bayesian inference and Bayes rule. Unfortunately, the

Here are the contents of notebooks/Chapter1_Introduction--16.md:

##Probability Distributions

Let's quickly recall what a probability distribution **is**: Let Z be some random vari

We can divide random variables into three classifications:

- **Z is discrete**: Discrete random variables may only assume values on a specified
- **Z is continuous**: Continuous random variable can take on arbitrarily exact valu
- **Z is mixed**: Mixed random variables assign probabilities to both discrete and con

###Discrete Case

If Z is discrete, then its distribution is called a *probability mass function*, which

$$P(Z = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k=0,1,2, \dots$$

What is λ ? It is called the parameter, and it describes the shape of the distrib

Unlike λ , which can be any positive number, the value k in the above formula m

If a random variable Z has a Poisson mass distribution, we denote **this** by writing

$$Z \sim \text{Poi}(\lambda)$$

One very useful property of the Poisson random variable, given we know λ , is tha

$$E\left[Z \mid \lambda \right] = \lambda$$

We will use **this** property often, so it's something useful to remember. Below we plot the

Here are the contents of notebooks/Chapter1_Introduction--17-input.py:

```
figsize( 12.5, 4)
```

```
import scipy.stats as stats
```

```

a = np.arange( 16 )
poi = stats.poisson
lambda_ = [1.5, 4.25 ]

plt.bar( a, poi.pmf( a, lambda_[0]), color=colours[0],
        label = "$\lambda = %.1f$" % lambda_[0], alpha = 0.60,
        edgecolor = colours[0], lw = "3")

plt.bar( a, poi.pmf( a, lambda_[1]), color=colours[1],
        label = "$\lambda = %.1f$" % lambda_[1], alpha = 0.60,
        edgecolor = colours[1], lw = "3")

plt.xticks( a + 0.4, a )
plt.legend()
plt.ylabel("probability of $k$")
plt.xlabel("$k$")
plt.title("Probability mass function of a Poisson random variable; differing \
$\lambda$ values");

```

Here are the contents of notebooks/Chapter1_Introduction--17-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--18.md:

###Continuous Case

Instead of a probability mass function, a continuous random variable has a *probability

$f_Z(z | \lambda) = \lambda e^{-\lambda z}$, $z \geq 0$

Like the Poisson random variable, an exponential random variable can only take on non-negative values.

When a random variable Z has an exponential distribution with parameter λ , we

$Z \sim \text{Exp}(\lambda)$

Given a specific λ , the expected value of an exponential random variable is equal to

$E[Z] = \frac{1}{\lambda}$

Here are the contents of notebooks/Chapter1_Introduction--19-input.py:

```

a = np.linspace(0,4, 100)
expo = stats.expon
lambda_ = [0.5, 1]

for l,c in zip(lambda_,colours):
    plt.plot( a, expo.pdf( a, scale=1./l), lw=3,

```



```

        color=c, label = "$\lambda = %.1f$" % l)
plt.fill_between( a, expo.pdf( a, scale=1./l), color=c, alpha = .33)

plt.legend()
plt.ylabel("PDF at $z$")
plt.xlabel("$z$")
plt.title("Probability density function of an Exponential random variable;\
differing $\lambda$");

```

Here are the contents of notebooks/Chapter1_Introduction--19-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--20.md:

```
###But what is  $\lambda$  ;$?
```

****This question is what motivates statistics**.** In the real world, λ is hidden from us.

Bayesian inference is concerned with **beliefs** about what λ is. Rather than try to estimate it, we try to model our beliefs about it.

This might seem odd at first: after all, λ is fixed, it is not (necessarily) random.

Here are the contents of notebooks/Chapter1_Introduction--21.md:

```
##### Example: Inferring behaviour from text-message data
```

Let's try to model a more interesting example, concerning text-message rates:

> You are given a series of text-message counts from a user of your system. The data, plotted as a bar chart, looks like this:

Here are the contents of notebooks/Chapter1_Introduction--22-input.py:

```

figsize( 12.5, 3.5 )
count_data = np.loadtxt("data/txtdata.csv")
n_count_data = len(count_data)
plt.bar( np.arange( n_count_data ), count_data, color = "#348ABD" )
plt.xlabel( "Time (days)" )
plt.ylabel("count of text-msgs received")
plt.title("Did the user's texting habits change over time?")
plt.xlim( 0, n_count_data );

```

Here are the contents of notebooks/Chapter1_Introduction--22-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--23.md:

Before we begin, with respect to the plot above, would you say there was a change in behaviour during the time period?

How can we start to model *this*? Well, *as* I conveniently already introduced, a Poisson random variable

```
$$ C_i \sim \text{Poisson}(\lambda) $$
```

We are not sure about what the λ parameter *is* though. Looking at the chart above

How can we mathematically represent *this*? We can think, that at some later date (call it τ)

```
$$
\lambda =
\begin{cases}
\lambda_1 & \text{if } t < \tau \\
\lambda_2 & \text{if } t \geq \tau
\end{cases}
$$
```

If, *in* reality, no sudden change occurred and indeed $\lambda_1 = \lambda_2$, the λ parameter is

We are interested *in* inferring the unknown λ s. To use Bayesian inference, we need to specify priors

```
\begin{align}
&\lambda_1 \sim \text{Exp}(\alpha) \\
&\lambda_2 \sim \text{Exp}(\alpha)
\end{align}
```

α *is* called a *hyper-parameter*, or a *parent-variable*, literally a parameter that defines the parameters

```
$$ \frac{1}{N} \sum_{i=0}^N C_i \approx E[\lambda; \alpha] = \frac{1}{\alpha} $$
```

Alternatively, and something I encourage the reader to *try*, *is* to have two priors: one for λ_1 and one for λ_2

What about τ ? Well, due to the randomness, it *is* too difficult to pick out when the change occurred

```
\begin{align}
&\tau \sim \text{DiscreteUniform}(1, 70) \\
&\rightarrow P(\tau = k) = \frac{1}{70}
\end{align}
```

So after all *this*, what does our overall prior *for* the unknown variables look like? Frankly, it's messy

Introducing our first hammer: PyMC

```
-----
```

PyMC **is** a Python library **for** programming Bayesian analysis [3]. It **is** a fast, well-maintained library. We will model the above problem using the PyMC library. This type of programming **is** called probabilistic programming. B. Cronin [5] has a very motivating description of probabilistic programming:

> Another way of thinking about **this**: unlike a traditional program, which only runs **in** memory, probabilistic programs are executed by sampling from a joint distribution. Due to its poorly understood title, I'll refrain from using the name *probabilistic programming*. The PyMC code **is** easy to follow along: the only novel thing should be the syntax, and I will try to keep it as simple as possible.

Here are the contents of notebooks/Chapter1_Introduction--24-input.py:

```
import pymc as mc

alpha = 1.0/count_data.mean() #recall count_data is
                               #the variable that holds our txt counts

lambda_1 = mc.Exponential( "lambda_1", alpha )
lambda_2 = mc.Exponential( "lambda_2", alpha )

tau = mc.DiscreteUniform( "tau", lower = 0, upper = n_count_data )
```

Here are the contents of notebooks/Chapter1_Introduction--25.md:

In the above code, we create the PyMC variables corresponding to λ_1 , λ_2 , and τ .

Here are the contents of notebooks/Chapter1_Introduction--26-input.py:

```
print "Random output:", tau.random(),tau.random(), tau.random()
```

Here are the contents of notebooks/Chapter1_Introduction--26-output-0.txt:

Here are the contents of notebooks/Chapter1_Introduction--27-input.py:

```
@mc.deterministic
def lambda_( tau = tau, lambda_1 = lambda_1, lambda_2 = lambda_2 ):
    out = np.zeros( n_count_data )
    out[:tau] = lambda_1 #lambda before tau is lambda1
    out[tau:] = lambda_2 #lambda after tau is lambda2
    return out
```

Here are the contents of notebooks/Chapter1_Introduction--28.md:

This code **is** creating a new **function** `lambda_`, but really we think of it **as** a random variable.

Here are the contents of notebooks/Chapter1_Introduction--29-input.py:

```
observation = mc.Poisson( "obs", lambda_, value = count_data, observed = True)

model = mc.Model( [observation, lambda_1, lambda_2, tau] )
```

Here are the contents of notebooks/Chapter1_Introduction--30.md:

The variable `observation` combines our data, `count_data`, with our proposed data-gener

The below code will be explained in the Chapter 3, but `this is` where our results come fr

Here are the contents of notebooks/Chapter1_Introduction--31-input.py:

```
### Mysterious code to be explained in Chapter 3.
mcmc = mc.MCMC(model)
mcmc.sample( 40000, 10000, 1 )
```

Here are the contents of notebooks/Chapter1_Introduction--31-output-0.txt:

Here are the contents of notebooks/Chapter1_Introduction--31-output-1.txt:

Here are the contents of notebooks/Chapter1_Introduction--32-input.py:

```
lambda_1_samples = mcmc.trace( 'lambda_1' )[:]
lambda_2_samples = mcmc.trace( 'lambda_2' )[:]
tau_samples = mcmc.trace( 'tau' )[:]
```

Here are the contents of notebooks/Chapter1_Introduction--33-input.py:

```
figsize(12.5, 10)
#histogram of the samples:

ax = plt.subplot(311)
ax.set_autoscaley_on(False)

plt.hist( lambda_1_samples, histtype='stepfilled', bins = 30, alpha = 0.85,
          label = "posterior of  $\lambda_1$ ", color = "#A60628",normed = True )
plt.legend(loc = "upper left")
plt.title(r"Posterior distributions of the variables  $\lambda_1, \lambda_2, \tau$ ")
plt.xlim([15,30])
plt.xlabel(" $\lambda_2$  value")
plt.ylabel("probability")

ax = plt.subplot(312)
ax.set_autoscaley_on(False)

plt.hist( lambda_2_samples,histtype='stepfilled', bins = 30, alpha = 0.85,
```

```

        label = "posterior of  $\lambda_2$ ", color="#7A68A6", normed = True )
plt.legend(loc = "upper left")
plt.xlim([15,30])
plt.xlabel(" $\lambda_2$  value")
plt.ylabel("probability")

plt.subplot(313)

w = 1.0/ tau_samples.shape[0] * np.ones_like( tau_samples )
plt.hist( tau_samples, bins = n_count_data, alpha = 1,
        label = r"posterior of  $\tau$ ",
        color="#467821", weights=w, rwidth =2. )
plt.xticks( np.arange( n_count_data ) )

plt.legend(loc = "upper left");
plt.ylim([0,.75])
plt.xlim([35, len(count_data)-20])
plt.xlabel(" $\tau$  (in days)")
plt.ylabel("probability");

```

Here are the contents of notebooks/Chapter1_Introduction--33-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--34.md:

Interpretation

Recall that the Bayesian methodology returns a *distribution*, hence we now have distrib

Also notice that the posterior distributions for the λ 's do not look like any ex

Our analysis also returned a distribution for what τ might be. Its posterior distri

Here are the contents of notebooks/Chapter1_Introduction--35.md:

Why would I want samples from the posterior, anyways?

We will deal with this question for the remainder of the book, and it is an understatement

In the code below, we are calculating the following: Let i index samples from the post

Here are the contents of notebooks/Chapter1_Introduction--36-input.py:

```

figsize( 12.5, 5)
# tau_samples, lambda_1_samples, lambda_2_samples contain
# N samples from the corresponding posterior distribution

```

```

N = tau_samples.shape[0]
expected_texts_per_day = np.zeros(n_count_data)
for day in range(0, n_count_data):
    # ix is a bool index of all tau samples corresponding to
    # the switchpoint occurring prior to value of 'day'
    ix = day < tau_samples
    # Each posterior sample corresponds to a value for tau.
    # for each day, that value of tau indicates whether we're "before"
    # (in the lambda1 "regime") or
    # "after" (in the lambda2 "regime") the switchpoint.
    # by taking the posterior sample of lambda1/2 accordingly, we can average
    # over all samples to get an expected value for lambda on that day.
    # As explained, the "message count" random variable is Poisson distributed,
    # and therefore lambda (the poisson parameter) is the expected value of "message count"
    expected_texts_per_day[day] = (lambda_1_samples[ix].sum()
                                   + lambda_2_samples[~ix].sum() ) / N

plt.plot( range( n_count_data), expected_texts_per_day, lw =4, color = "#E24A33",
          label = "expected number of text-messages recieved")
plt.xlim( 0, n_count_data )
plt.xlabel( "Day" )
plt.ylabel( "Expected # text-messages" )
plt.title( "Expected number of text-messages received")
plt.ylim( 0, 50 )
plt.bar( np.arange( len(count_data) ), count_data, color = "#348ABD", alpha = 0.65,
         label="observed texts per day")

plt.legend(loc="upper left");

```

Here are the contents of notebooks/Chapter1_Introduction--36-output-0.png:

Here are the contents of notebooks/Chapter1_Introduction--37.md:

Our analysis shows strong support for believing the user's behavior did change (λ).

Here are the contents of notebooks/Chapter1_Introduction--38.md:

Exercises

1\). Using `lambda_1_samples` and `lambda_2_samples`, what is the mean of the posterior

Here are the contents of notebooks/Chapter1_Introduction--39-input.py:

#type your code here.

Here are the contents of notebooks/Chapter1_Introduction--40.md:

2\). What **is** the expected percentage increase **in** text-message rates? **hint:** compute the

Here are the contents of notebooks/Chapter1_Introduction--41-input.py:

#type your code here.

Here are the contents of notebooks/Chapter1_Introduction--42.md:

3\). What **is** the mean of λ_1 ****given**** we know τ **is** less than **45**. That **is**,

Here are the contents of notebooks/Chapter1_Introduction--43-input.py:

#type your code here.

Here are the contents of notebooks/Chapter1_Introduction--44.md:

References

- [1] Gelman, Andrew. N.p.. Web. 22 Jan 2013. <http://andrewgelman.com/2005/07/n_is_nev
- [2] Norvig, Peter. 2009. [*The Unreasonable Effectiveness of Data*] (<http://www.csee.w>
- [3] Patil, A., D. Huard and C.J. Fonnesbeck. 2010. PyMC: Bayesian Stochastic Modelling **in** Python. Journal of Statistical Software, 35(4), pp. 1-81.
- [4] Jimmy Lin and Alek Kolcz. Large-Scale Machine Learning at Twitter. Proceedings of
- [5] Cronin, Beau. "Why Probabilistic Programming Matters." 24 Mar 2013. Google, Online

Here are the contents of notebooks/Chapter1_Introduction--45-input.py:

```
from IPython.core.display import HTML
def css_styling():
    styles = open("../styles/custom.css", "r").read()
    return HTML(styles)
css_styling()
```

Here are the contents of notebooks/Chapter1_Introduction--46-input.py: